

Automatic differentiation beyond typedef and operator overloading

Peter Caspers

Quaternion Risk Management

November 30, 2015

Table of contents

- 1 Introduction to AD
- 2 Approaches in QuantLib
- 3 Source code transformation
- 4 Questions

AD in a nutshell 1/2

- compute derivatives by symbolic differentiation of an operation sequence ($+$ $*$ $-$ $/$, \exp , \sin , ...) using the chain rule
- results are exact up to machine precision, also in higher orders
- implementation:
 - operator overloading tools instrumenting the double type (e.g. CppAD, ADOL-C, Adept, dco, proprietary tools)
 - source code transformation tools (e.g. ADIC, OpenAD/F)
 - coding by hand

AD in a nutshell 2/2

- one forward sweep yields one directional derivative of the vector of output variables
- one reverse sweep yields the gradient w.r.t. all input variables of a linear combination of the output variables
- the complexity for one (forward or reverse) sweep is a constant, low multiple of the complexity for one function evaluation¹
- in particular: law of cheap gradient !

¹theory claims that the multiple in adjoint mode is bounded by 4

The typedef approach

- just says `typedef CppAD::AD<double> Real`
- it is a bit more complicated than that
- QuantLibAdjoint (CompatibL / Alex Sokol), with additional logic (tapescript)
- *AD-or-not-AD* decision at compile time and globally, i.e. no selective activation of variables

Matrix multiplication with (sleeping) active doubles

```
Matrix_t<T> A(1024, 1024);
Matrix_t<T> B(1024, 1024);
...
Matrix_t<T> C = A * B;
```

- T = double: 764 ms
- T = CppAD::AD<double>: 8960 ms
- Penalty: $11.7x^2$
- the compiler does not seem to apply certain optimizations to the active type that are possible for the native double (SIMD instructions)
- this is not an exception, but seems to occur for every “numerically intense” code section (see below for a second example)

²different OO tools may show different penalties; for a MinimalWrapper consisting of a double and a native MinimalWrapper - pointer (which is null always), the penalty is around 2.0x

The template approach

- introduce templated versions of relevant classes (e.g. `Matrix_t`)
- for backward compatibility, `typedef Matrix_t<Real> Matrix`
- it is a bit more complicated than that
- allows mixing of active and native classes, as required, i.e. activation of variables in selected parts of the application only
- work in progress, but basic IRD stuff works (like yield and volatility termstructures, swaps, CMS coupons, GSR model)
- <https://github.com/pcaspers/quantlib/tree/adjoint>
- <https://quantlib.wordpress.com/tag/automatic-differentiation/>

Expensive gradients with operator overloading

- the typedef as well as the template approach use operator overloading tools (like CppAD)
- for numerically intense algorithms, we observe dramatic performance loss (because less optimization can be applied to non-native types)
- e.g. a convolution engine for bermudan swaptions is 80x slower³ in adjoint mode compared to one native-double pricing
- if AD is actually not needed, the template approach is the way out, otherwise we need other techniques

³see <https://quantlib.wordpress.com/2015/04/14/adjoint-greeks-iv-exotics>

Source Code Transformation

- generate adjoint code at compile time, which may yield better performance
- however, does not work out of the box like OO tools
- no mature tool for C++ (ADIC 2.0 = “OpenAD/Cpp” under development)
- needs specific preparation of code before it can be applied


OpenAD/F

- OpenAD is a language independent AD backend working with abstract xml representations (XAIF) of the computational model
- OpenAD/F adds a Fortran 90 front end
- Open Source, proven on large scale real-world models
- <http://www.mcs.anl.gov/OpenAD>

From QuantLib to SCT

- isolate the core computational code and reimplement it in Fortran
- use OpenAD/F to generate adjoint code, build a separate support library from that
- use a wrapper class on the QuantLib side to communicate with the support library
- minimal library example ⁴ and LGM swaption engine⁵ available
- build via make (AD support library) or make plain (without OpenAD - transformation, for testing)

⁴ <https://github.com/pcaspers/quantlib/tree/master/QuantLibOAD/simplelib>

⁵ <https://github.com/pcaspers/quantlib/tree/master/QuantLibOAD/lgm> 

LGM Bermudan swaption convolution engine

- core computation can be implemented in around 200 lines
- native interface only using doubles and arrays of doubles
- input: relevant times $\{t_i\}$, model $\{(H(t_i), \zeta(t_i), P(0, t_i))\}$, Termsheet, codified as index lists $\{k_i, l_i, \dots\}$
- output: npv, gradient w.r.t. $\{(H(t_i), \zeta(t_i), P(0, t_i))\}$

```
subroutine lgm_swaption_engine(n_times, times, modpar, n_expiries, &
    expiries, callput, n_floats, &
    float_startidxs, float_mults, index_acctimes, float_spreads, &
    float_t1s, float_t2s, float_tps, &
    fix_startidxs, n_fixs, fix_cpn, fix_tps, &
    integration_points, stddevs, res)
```

Building the AD support library

```

emacs@peter-ThinkPad-W520
File Edit Options Buffers Tools Compile Help
[*- mode: compilation; default-directory: "~/quantlib/QuantLib0AD/lgm/" -*-
Compilation started at Sat Oct 31 09:52:42

make
openad -c -m rj lgm.f90
openad log: openad.2015-10-31_09:52:42.log-
preprocessing fortran
parsing preprocessed fortran
analyzing source code and translating to xaif
adjoint transformation
  getting runtime support file OAD_active.f90
  getting runtime support file w2f__types.f90
  getting runtime support file iaddr.c
  getting runtime support file ad_inline.f
  getting runtime support file OAD_cp.f90
  getting runtime support file OAD_rev.f90
  getting runtime support file OAD_tape.f90
  getting template file
translating transformed xaif to whirl
unparsing transformed whirl to fortran
postprocessing transformed fortran
gfortran -g -O3 -o OAD_active.o -c OAD_active.f90 -fpic
gfortran -g -O3 -o OAD_cp.o -c OAD_cp.f90 -fpic
gfortran -g -O3 -o OAD_tape.o -c OAD_tape.f90 -fpic
gfortran -g -O3 -o OAD_rev.o -c OAD_rev.f90 -fpic
gfortran -g -O3 -o driver_lgm.o -c driver_lgm.f90 -fpic
gfortran -g -O3 -o lgm.pre.xb.x2w.w2f.post.o -c lgm.pre.xb.x2w.w2f.post.f90 -fpic
gfortran -shared -g -O3 -o liblgm.so w2f__types.o OAD_active.o OAD_cp.o OAD_tape.o OAD_rev.o driver_lgm.o lgm.pre.xb.x2w.w2f.post.o

Compilation finished at Sat Oct 31 09:52:50

U:~%*- *compilation* All L1 (Compilation:exit (0))
Compilation finished

```

LGM Bermudan swaption convolution engine

- C++ wrapper is a normal QuantLib pricing engine
- precomputes the values and organizes them in arrays for the Fortran core
- invokes the Fortran routine
- stores the npv and the adjoint gradient as results

```
void LgmSwaptionEngineAD::calculate() const {
    // collect data needed for core computation routine
    ...
    // join all dates and fill index vectors
    ...
    // call core computation routine and set results

    lgm_swaption_engine_ad_(&ntimes, &allTimes[0], &modpar[0], &nexpiries, ...
        &integration_pts, &std_devs, &res, &dres[0]);
    ...
    results_.value = res;
    results_.additionalResults["sensitivityTimes"] = allTimes;
    results_.additionalResults["sensitivityH"] = H_sensitivity;
    results_.additionalResults["sensitivityZeta"] = zeta_sensitivity;
    results_.additionalResults["sensitivityDiscount"] = discount_sensitivity;
```

Performance

- 10y bermudan swaption, yearly callable
- 49 grid points per expiry
- single pricing⁶ (non-transformed code): 4.2 ms
- pricing + gradient $\in \mathbb{R}^{105}$: 25.6 ms
- additional stuff⁷: 6.2 ms
- adjoint calculation multiple: 6.1x (7.6x including add. stuff)
- common, practical target for the adjoint multiple: 5x - 10x

⁶Intel(R) Core(TM) i7-2760QM CPU @ 2.40GHz, using one thread

⁷transformation of gradient w.r.t. model parameters to usual vegas, see below

How not to use AD

- avoid to record tapes that go through solvers, optimizers, etc.⁸
 - instead use the implicit function theorem to convert gradients w.r.t. calibrated (model) variables to gradients w.r.t. market variables
 - this is more efficient, less error prone (e.g. Bisection produces zero derivatives always, optimizations may produce bogus derivatives depending on the start value)
 - in the case of SCT as described above this is even necessary from a practical viewpoint
- apply AD only to differentiable programs (replace a digital payoff for example by a call spread)
- avoid to record *long* tapes (e.g. for *all* paths of a MC simulation), reuse a tape recorded (in a *tape-safe* way) on one path

⁸not to be confused with feeding AD - derivatives of the target function to optimizers like Levenberg-Marquardt or Newton-style solvers

Calibration of LGM model

To illustrate the usage of the implicit function theorem, consider the calibration to n swaptions⁹

$$\text{Black}(\sigma_1) - \text{Npv}_{\text{LGM}}(\zeta_1) = 0$$

...

$$\text{Black}(\sigma_n) - \text{Npv}_{\text{LGM}}(\zeta_n) = 0$$

with

$$\frac{\partial \text{Npv}_{\text{LGM}}}{\partial \zeta} = \text{diag}(\nu_1, \dots, \nu_n), \text{ all } \nu_i \neq 0 \quad (1)$$

⁹recall that $\zeta(t)$ is the accumulated model variance up to time t

Implicit function theorem

Locally, there exists a unique g

$$g(\sigma_1, \dots, \sigma_n) = (\zeta_1, \dots, \zeta_n) \quad (2)$$

and

$$\frac{\partial g}{\partial \sigma} = \left(\frac{\partial \text{Npv}_{\text{LGM}}}{\partial \zeta} \right)^{-1} \frac{\partial \text{Black}}{\partial \sigma} \quad (3)$$

Pasting the vega together

$$\frac{\partial \text{Npv}_{\text{Berm}}}{\partial \sigma} = \frac{\partial \text{Npv}_{\text{Berm}}}{\partial \zeta} \frac{\partial \zeta}{\partial \sigma} = \frac{\partial \text{Npv}_{\text{Berm}}}{\partial \zeta} \left(\frac{\partial \text{Npv}_{\text{Calib}}}{\partial \zeta} \right)^{-1} \frac{\partial \text{Black}}{\partial \sigma}$$

- the components can be calculated analytically (calibrating swaptions' market vegas) or using the ad engine (calibrating swaptions' ζ -gradient, but this is much cheaper than for the bermudan case)
- matrix inversion and multiplication is cheap
- the additional computation time is quite small (see the example above, the additional costs are the same as for 1.5x original NPV calculations)

Summary

- global instrumentation (via typedefs) with active variables can lead to performance (and memory) issues
- selective / mixed instrumentation (via templates) solves the issue, but leaves problems when AD is required for numerically intense parts of the code
- source code transformation can solve this issue, we gave an example in terms of a bermudan swaption engine transformed using OpenAD/F yielding an adjoint multiple of 6.1 compared to 80 with operator overloading (using CppAD)

Questions / Discussion

thank you for your attention